

Aryan Sanjay Patil

9344519501 | New York | ✉ arpatil@cs.stonybrook.edu |  [GitHub](#) |  [LinkedIn](#)

EDUCATION

Stony Brook University

Stony Brook, NY

Master of Science (MS) in Computer Science GPA: 3.85/4.0

Aug 2023 – May 2025

- Courses: Reinforcement Learning, Natural Language Processing, Data Science, Probability and Statistics.

University of Mumbai

Mumbai, India

Bachelor of Engineering (BE) in Computer Engineering GPA: 9.7/10.0

Aug 2019 – Aug 2023

- Courses: Data Structures, Computer Architecture, Operating Systems, Big Data Analytics, Cloud Computing.

SKILLS

Programming and tools: Python, C++, Java, SQL, R, C#, React, JavaScript, Git, Linux, Docker, Kubernetes

Frameworks: PyTorch, TensorFlow, Hugging Face, LangChain, OpenCV, ONNX, PySpark, CUDA, MLflow, FastAPI, Flask

AI Techniques: Transformers, VLLMs, VLMs, GANs, VAEs, RAGs, Diffusion Models, LoRA, Prompt Engineering, RL, MLOps

Cloud and Data: AWS (EC2, SageMaker, S3), Azure, GCP, MySQL, PostgreSQL, Pinecone, FAISS, Tableau, Power BI

EXPERIENCE

AI Engineer Intern

Mar 2025 – Present

SteadFast AI

Remote

- Contributed to LLM-based infrastructure and internal tooling for prompt workflows.
- Exploring and supporting the implementation of **RAG pipelines** within modular AI systems.

Research Project Assistant

Aug 2024 – Dec 2024

Stony Brook University

Stony Brook, NY

- Deployed a real-time emotion detection system for an e-learning platform on **AWS EC2** with **FastAPI**, using **multimodal RAG** to boost user engagement by **30%**.
- Enhanced CNN3D-ConvLSTM2D-based emotion recognition accuracy to **95%** with **23ms detection time** using **GAN**-augmented optical flow and linear algebra optimizations (**32% faster inference**).
- Integrated **SOTA** architectures (ViT, DeTR, UNet) for user tracking, improving real-world classification by 5%.

Machine Learning Research Assistant

Jun 2024 – Aug 2024

Brookhaven National Laboratory - Nuclear Physics | Publication in Progress

Upton, NY

- Led GPU-accelerated predictive modeling for Heavy-Ion Collisions on **3D data**, scaling inference across **NERSC** to cut analysis time from **1 hour to 2 seconds** and replace physical lab experiments.
- Developed a **Python module** using RandomForest Regressor, reducing error by **80%** (1e-4 to 2e-5) and optimizing model inference time by **50%** via **quantization** (FP32 to FP8), enabling efficient analysis.
- Finetuned 12 transformers with Gaussian embeddings for particle physics, reducing prediction noise to 1e-5 (**10x lower than RNNs**) and accelerating training by 40% via **CUDA** optimized **distributed clusters**.

Machine Learning Engineer Intern

Oct 2022 – May 2023

Dharmanandan Techno Projects Pvt Ltd

Mumbai, India

- Innovated a CNN3D variant with temporal attention for HAR in CCTV (**93% accuracy, 25% faster**), deployed via Docker on AWS using MLflow for model deployment, monitoring, and optimization (**50ms latency**).

PROJECTS

Reddit Analyzer App ([Live Demo](#)) | [Read Paper](#) | Multimodal, Boto3, Full-stack

- Built a real-time Reddit summarization app by fine-tuning **BART** and **Mistral 7B** and incorporated sarcasm detection (**LLaMA 7B**) to reduce hallucinations by **11%** and improve generalization.
- Formulated an end-to-end CI/CD NLP pipeline on **AWS SageMaker** (data prep, training, evaluation), reducing summarization time by **30%**, deployment time by **40%** and improving sentiment analysis by **10%**.

LocalityAI ([Live Demo](#)) | Recommendation System, Vector Database, React, Flask

- Prototyped a real-time ad recommendation engine using **GPT-4o** and multimodal prompts composed with a **text-to-image** workflow powered by a **QLoRA** fine-tuned **Stable Diffusion** model.
- Orchestrated a **RAG pipeline** (LangChain + FAISS) with hybrid search achieving a **25%** relevance boost.
- Reduced the latency of the system by 25% with **Redis** caching, running a modular Flask server on AWS EC2.

OPEN SOURCE CONTRIBUTIONS ([OSS PROFILE](#))

Published ML modules for LLM evaluation on **PyPi** with version control via Git, surpassing more than **3k downloads**.

Managed 32 repositories (Transformers, RL, ML) on [GitHub](#) with more than **100 stars**.

Merged 3 Pull Requests, including Preswald (2.5k stars), with work featured in the company's [blog](#).